

Page Rank Algorithm Using MPI

For this problem of assignment, I have implemented a Page Rank algorithm using **four different** approaches.

1. Using the Iterative Power method to find the eigenvector
2. Using MPI_Allreduce
3. Using MPI_Allgather
4. Using MPI_Irecv and MPI_Isend

I have also included all four methods codes with this report.

- **Iterative Power Method To Find Eigen Vector**

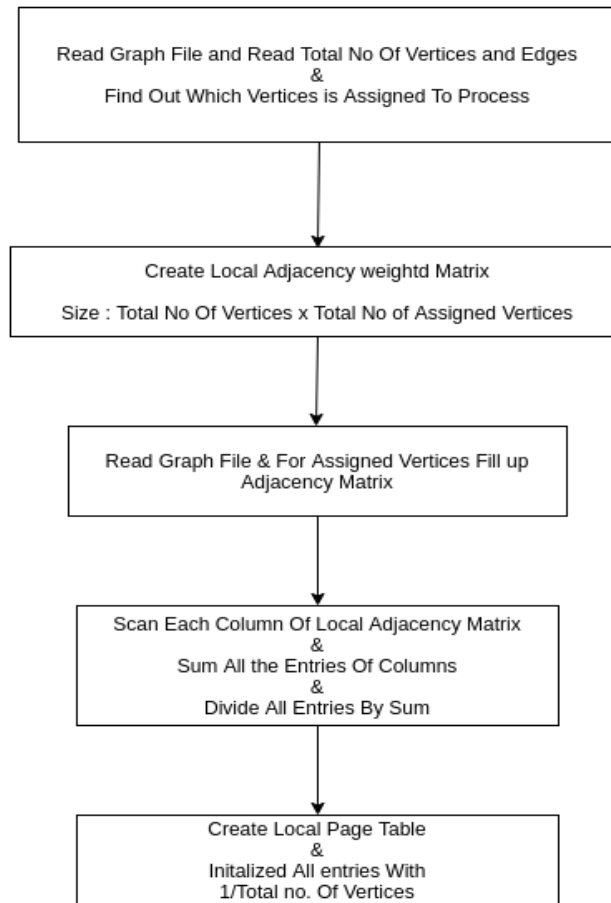
I have started my implementation with this method, In which we are creating transpose of a weighted adjacent matrix and iteratively we are multiplying a matrix with an initial guess of page rank till convergence.

This is an easy and straight forward method but after the implementation, I had come to the conclusion that If I will use adjacency matrix then it will be 50000 x 50000 for small graph and **Vector Matrix Multiplication of** (50000 x 50000) and (50000 x 1) will take humongous time.

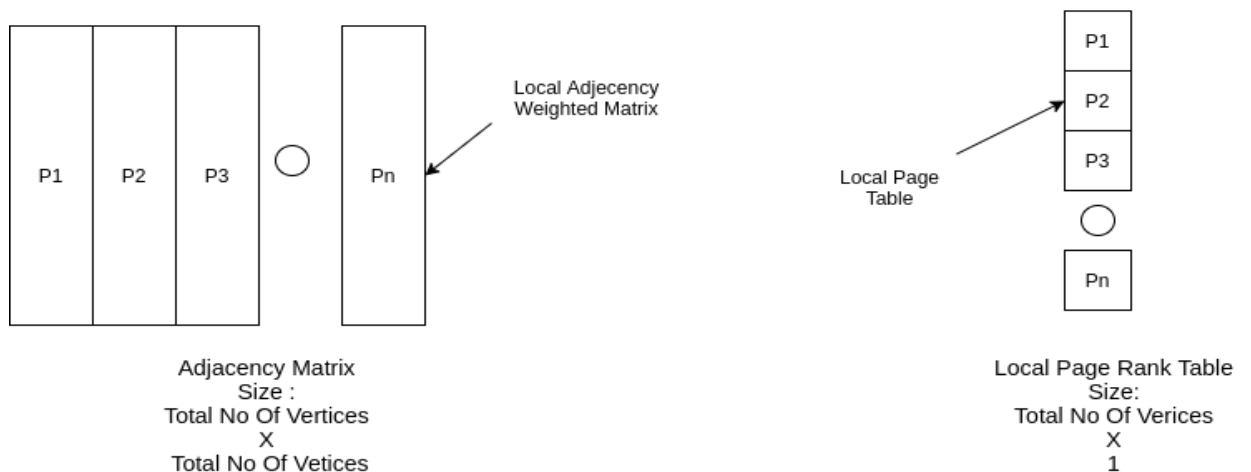
Even I was able to get output only for 64 processes and it was also taking too much time. In 1 Process even the program was not able to run on the Turing cluster.

Procedure :

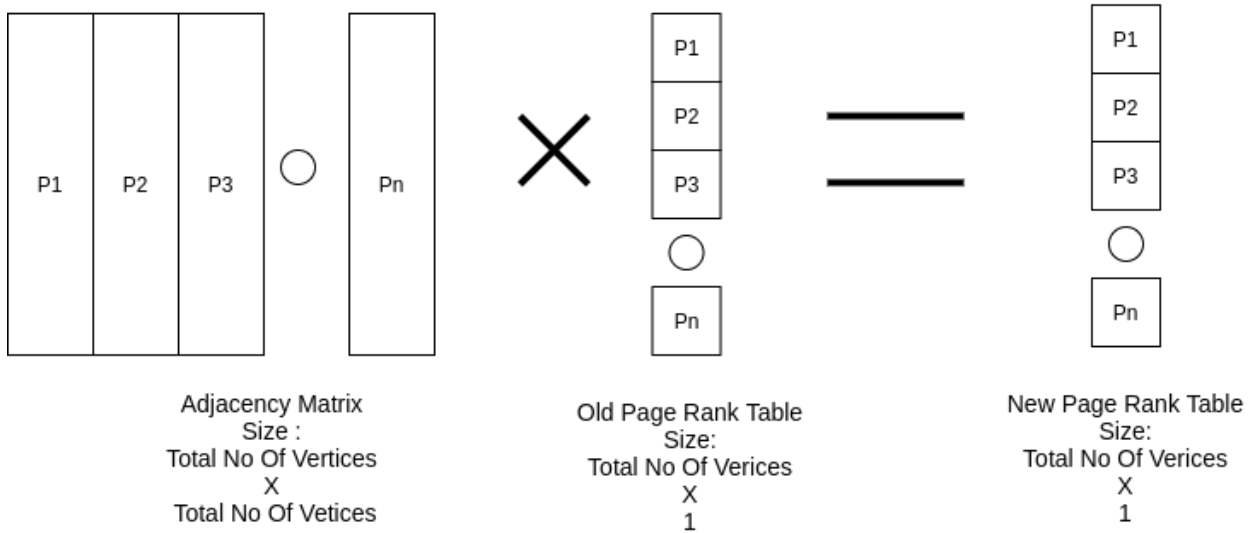
Procedure Followed By Each Process



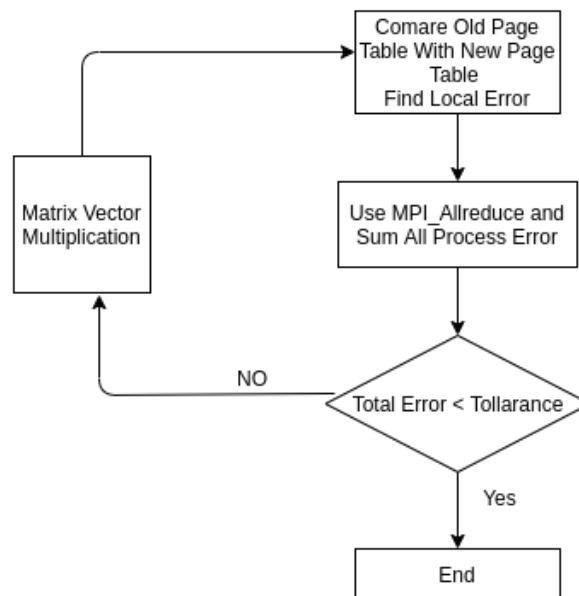
After The End Of Matrix And Local Page Table Creation Overall Picture



Matrix Vector Multiplication



After Each Matrix Multiplication



- **Using AllReduce Method :**

In this method, I adding the partial weight of each process to get the overall weight for the page rank table.

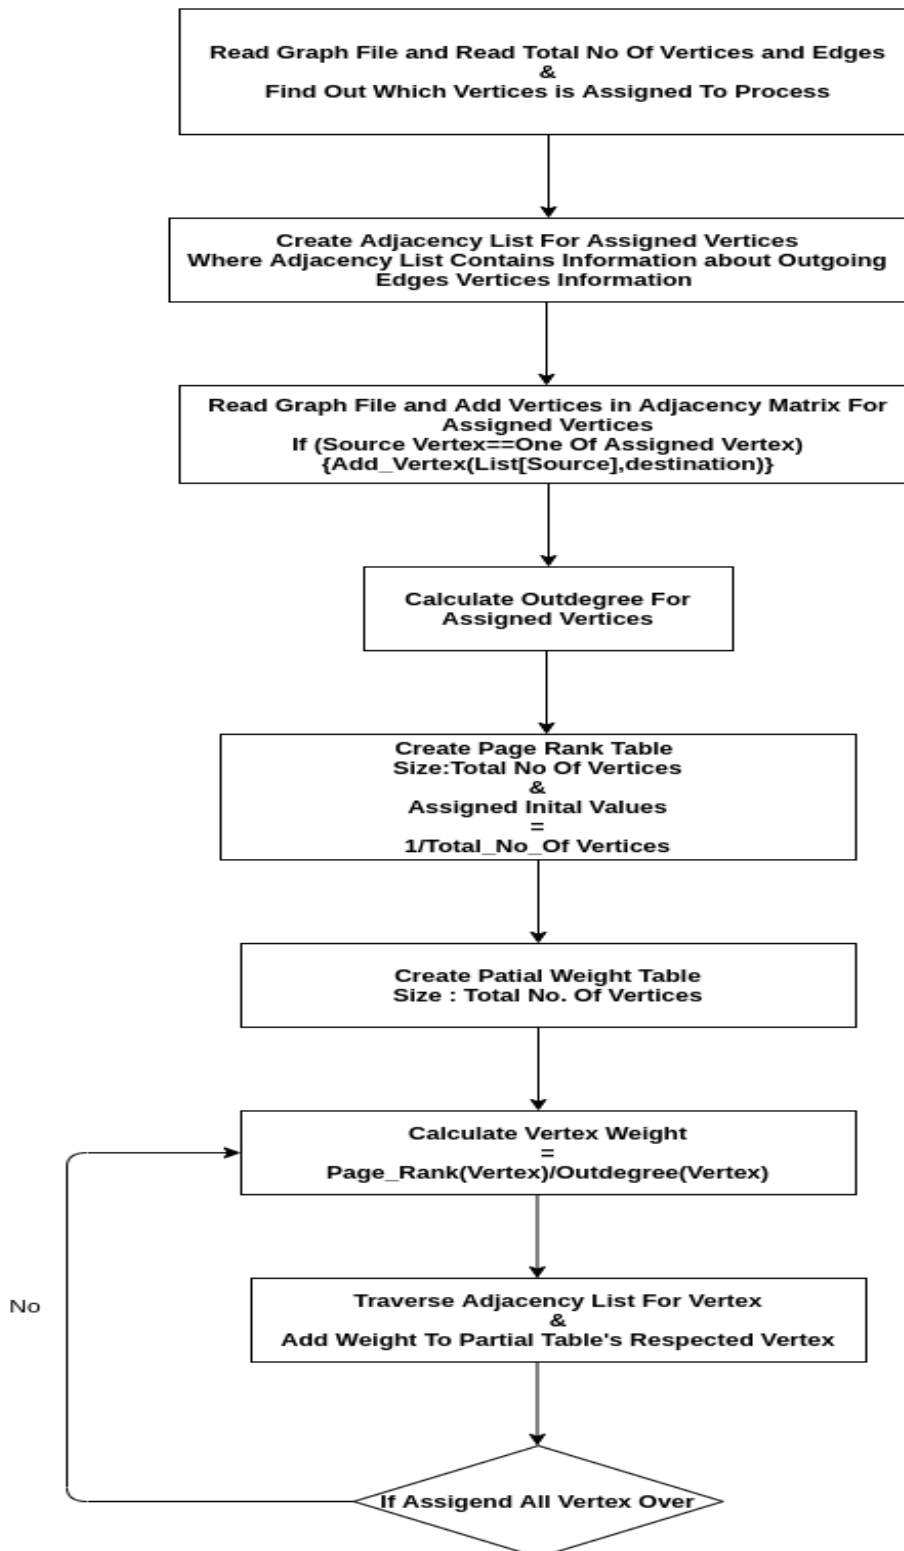
Here, I am getting a significant amount of performance improvement than the iterative method but as a no of process increase, my time is also increasing. Even sequential execution time is much lesser than parallel. So, I am getting no speed up here.

According to me, Not getting speed up to the sequential is,

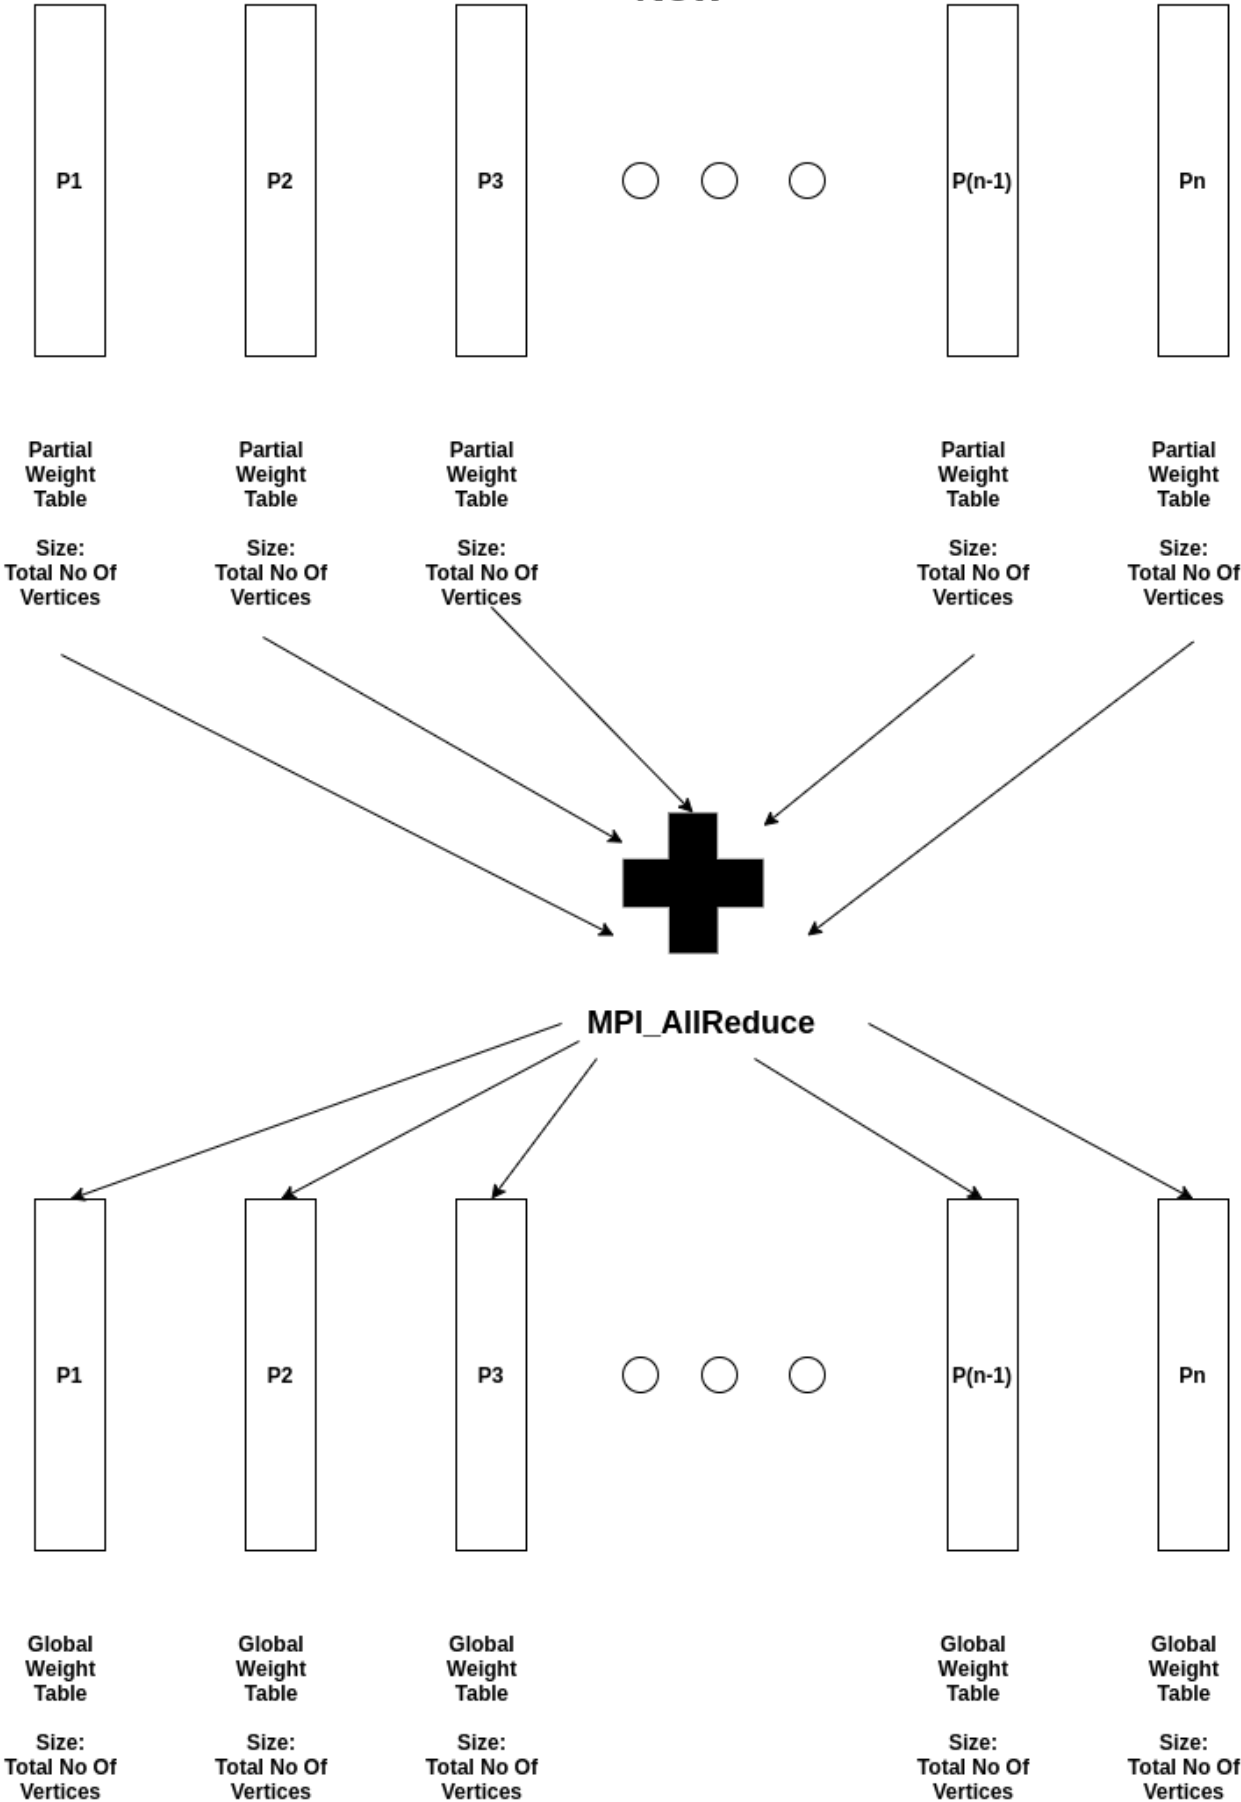
1. Communication cost is more than computation. Here we are gathering (50000 x number of process) size array over network then do the sum and again broadcast sum so I this communication will take a lot more time than computation which is not with a sequential program
2. MPI_Allreduce(). We know that MPI_Allreduce internally uses barrier before the beginning of Reduce operation so maybe it is slow down performance.

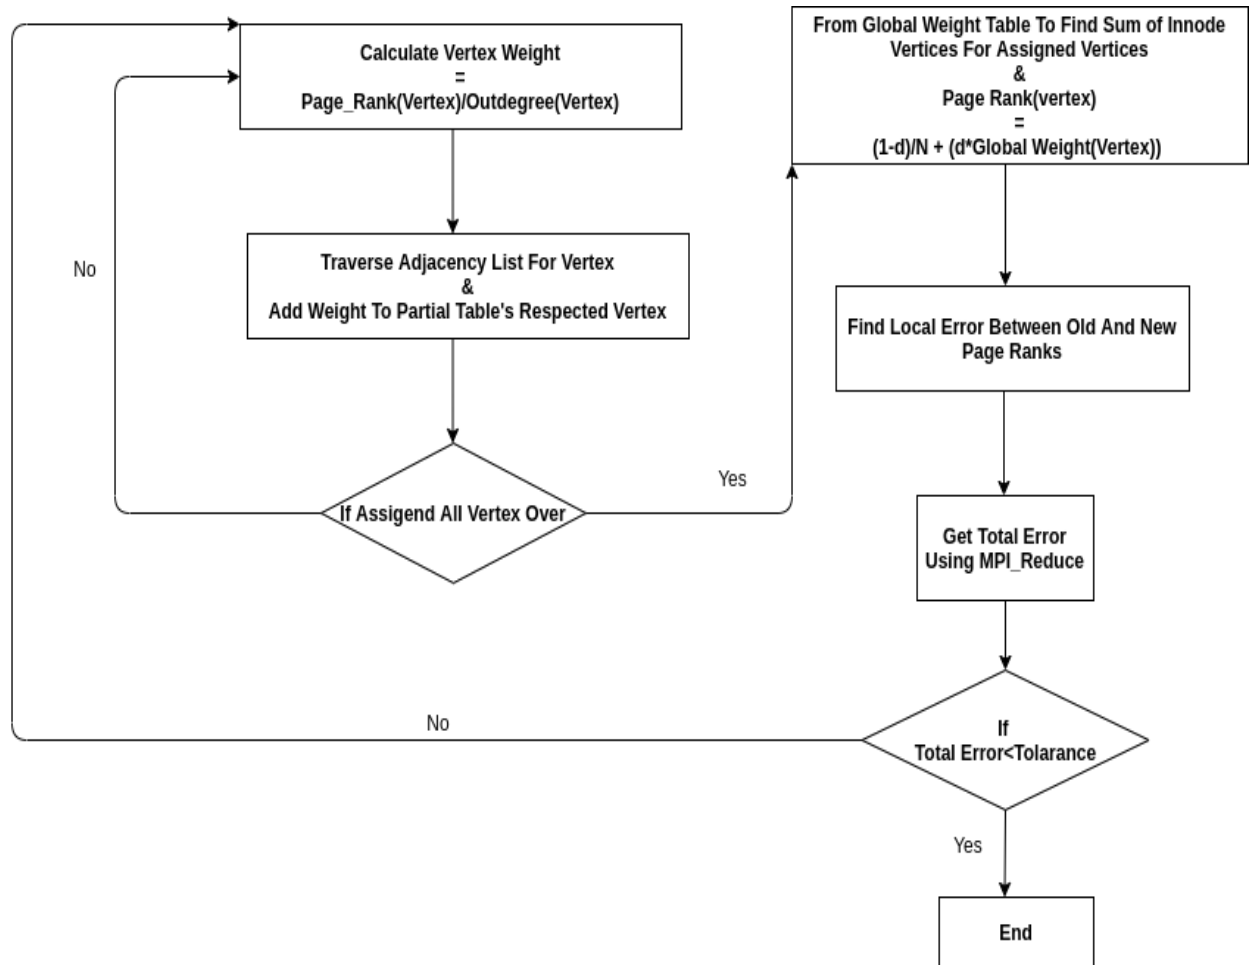
Process :

Procedure Followed By Each Process



Overall Picture After Completion Of Till
Now





- **Using MPI_Allgather:**

To get speed up, I think that, If I will reduce data that pass over the network then I will get up a speed up. So I come up with a solution and use MPI_Allgather.

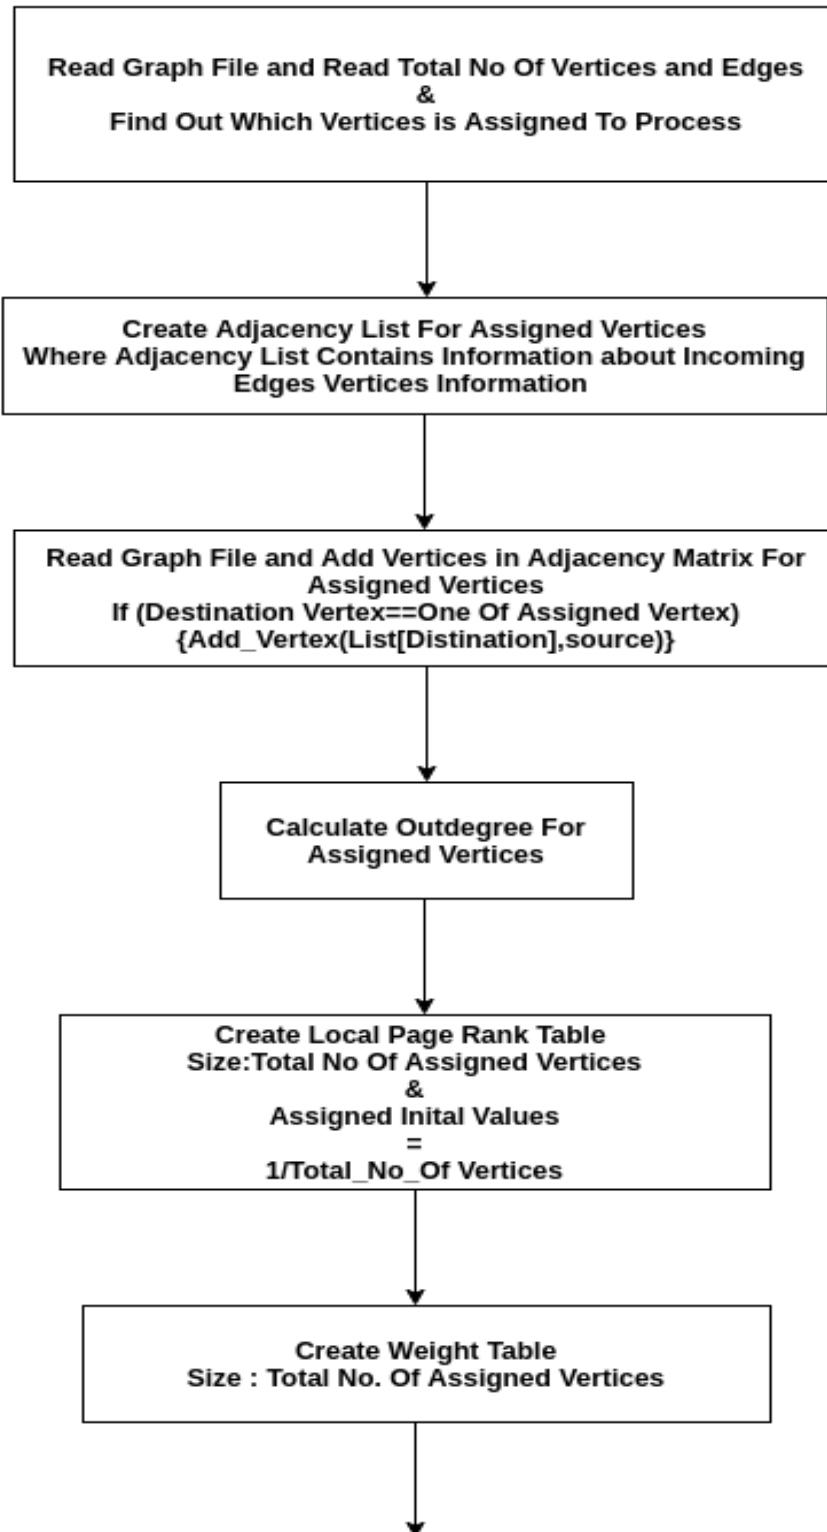
In this approach, I am only sending assigned vertices weight ($\text{PR}(\text{vertex}) / \text{Outdegree}(\text{vertex})$) to other processes and other processes use this weight to calculate a new pagerank.

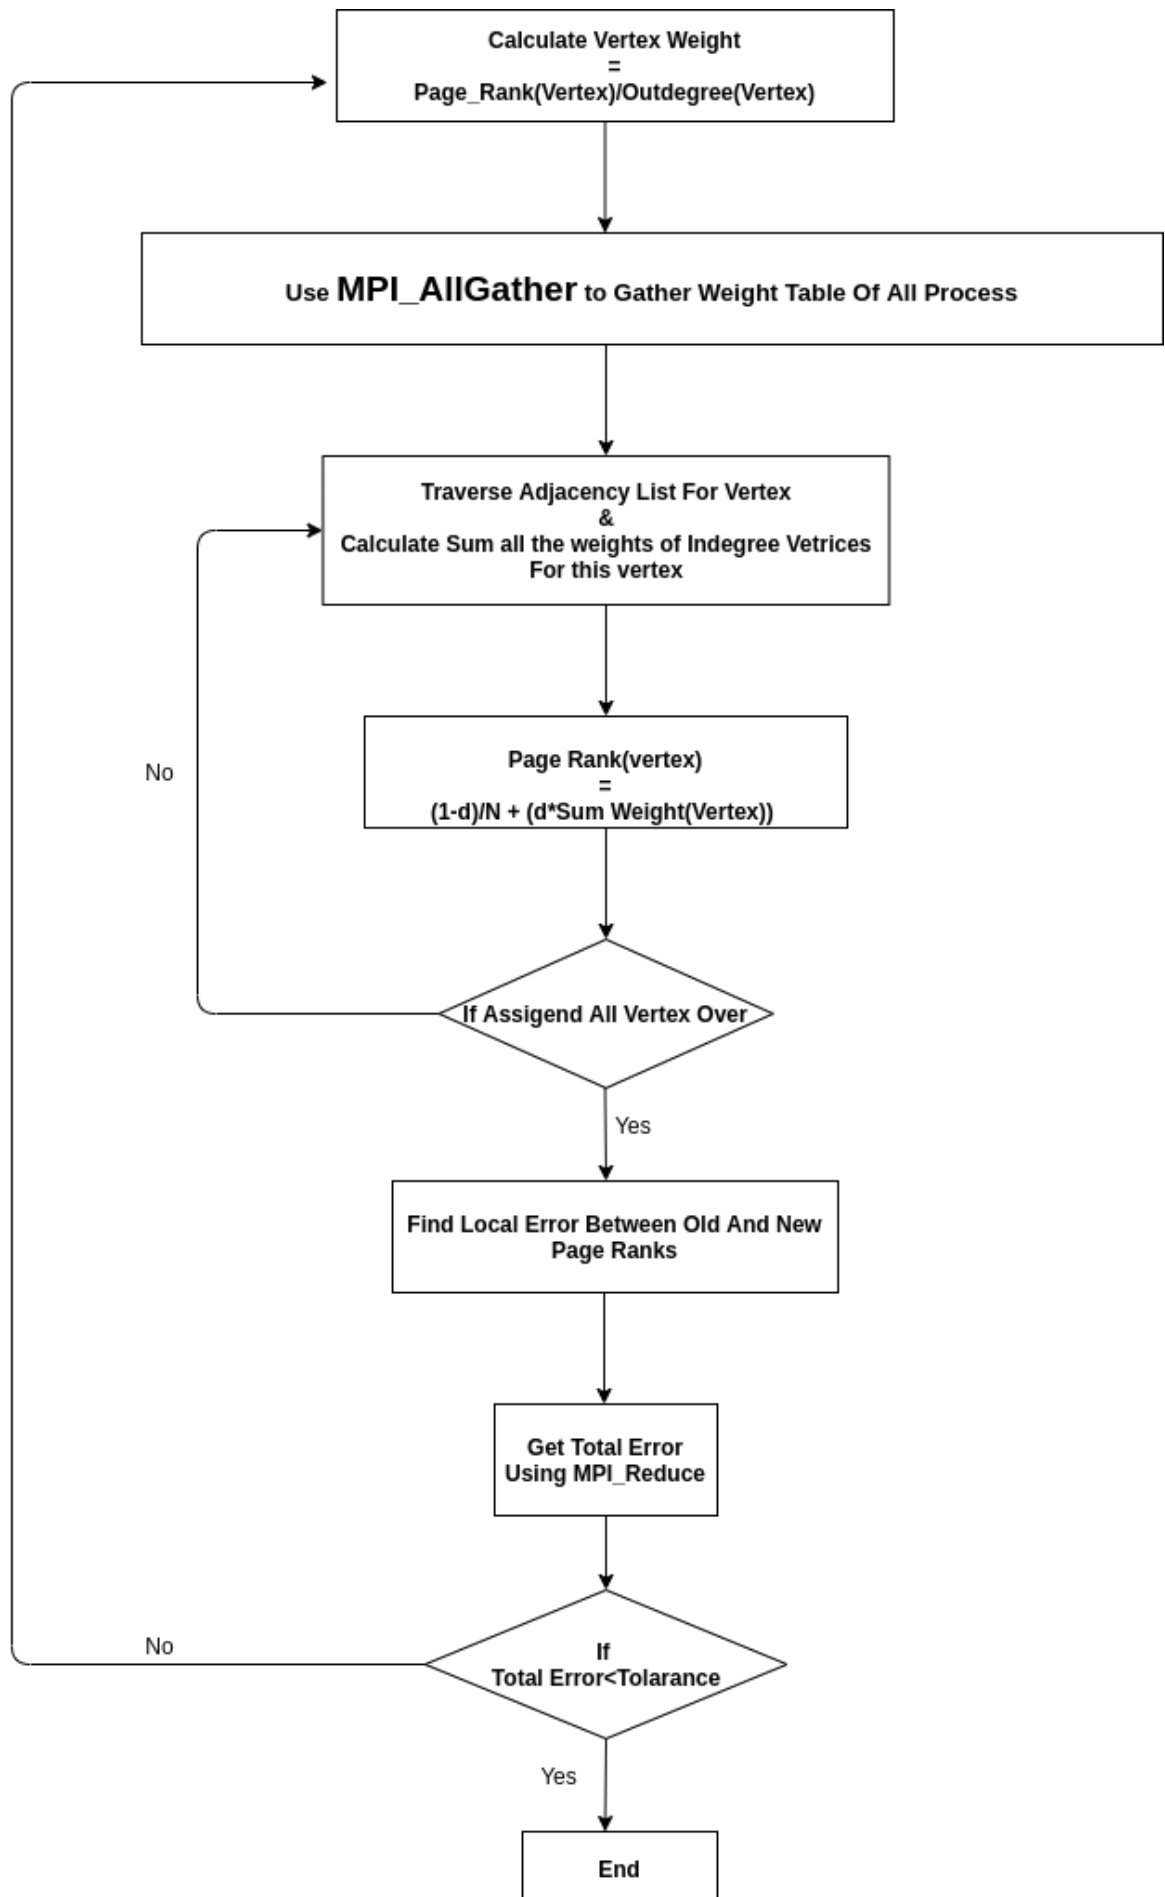
For a small graph, In MPI_Allreduce approach, we are sending (50000×1) vector while here We are sending $((50000)/\text{nproc}) \times 1$ vector.

So here I am getting slight improvement than the second approach but still, I am not getting speed up here.

Process :

Procedure Followed By Each Process





- **Using MPI_Isend And MPI_Irecv :**

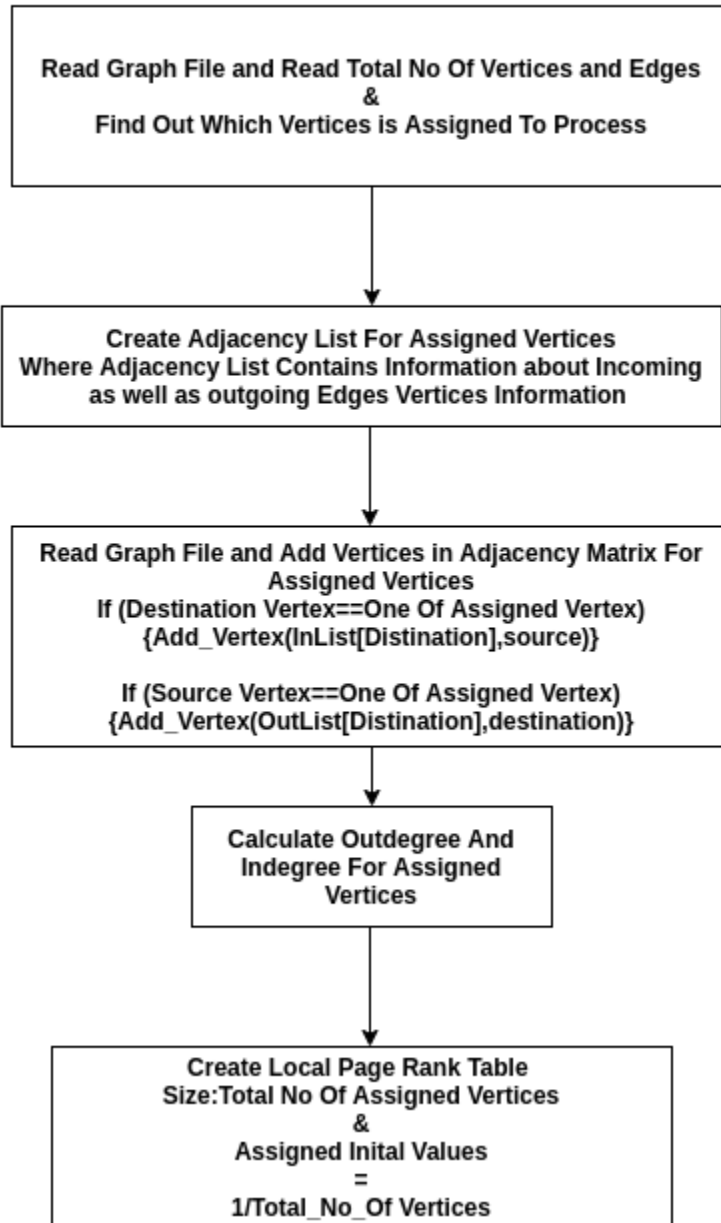
In all the previous two approaches, we are sending data to the other process whether it's required or not, and also with a blocking communication call.

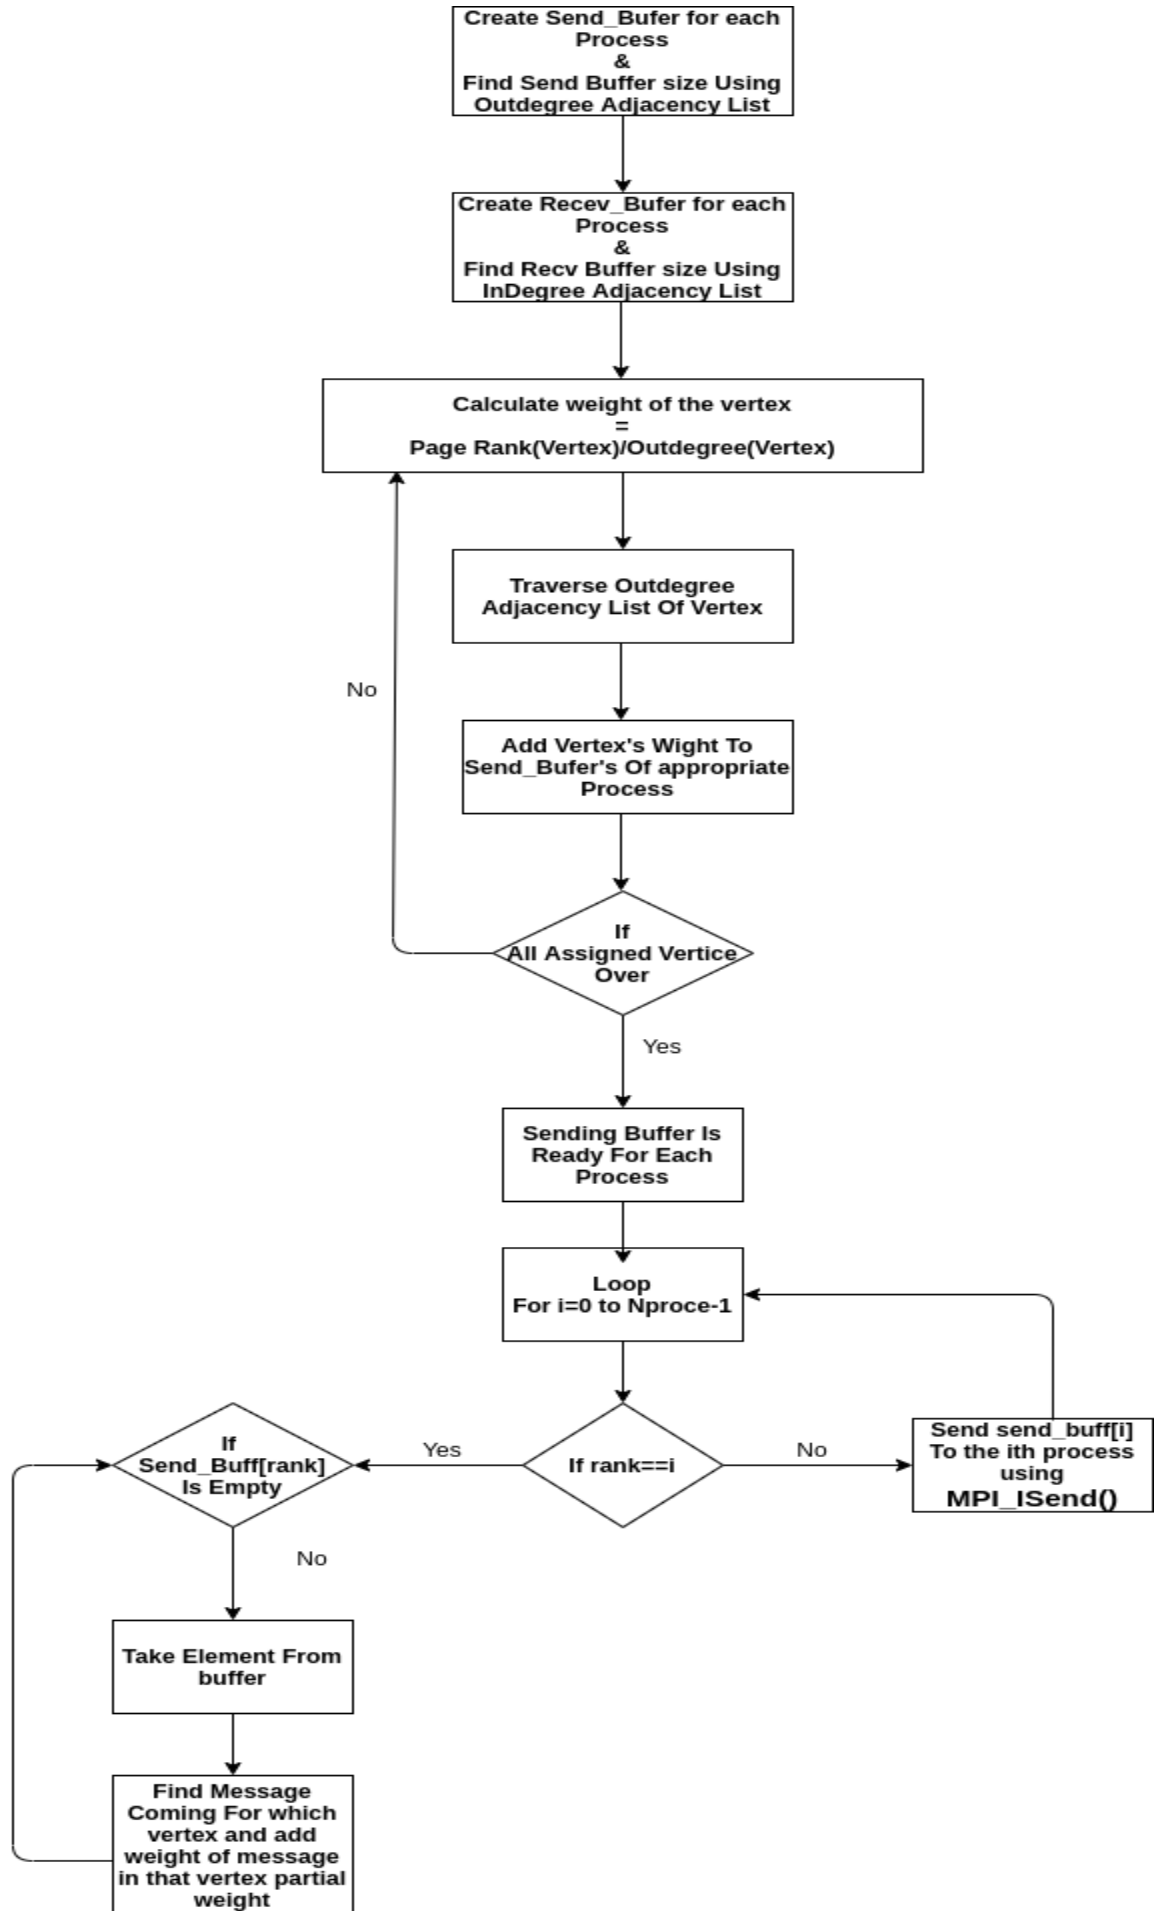
So, I get an idea that to just send weights to the process which requires only i.e communicate with the only boundary process. This will reduce lots of communication overhead and we can send this data using MPI_Isend and MPI_Irecv. Which will also overlap communication with computation.

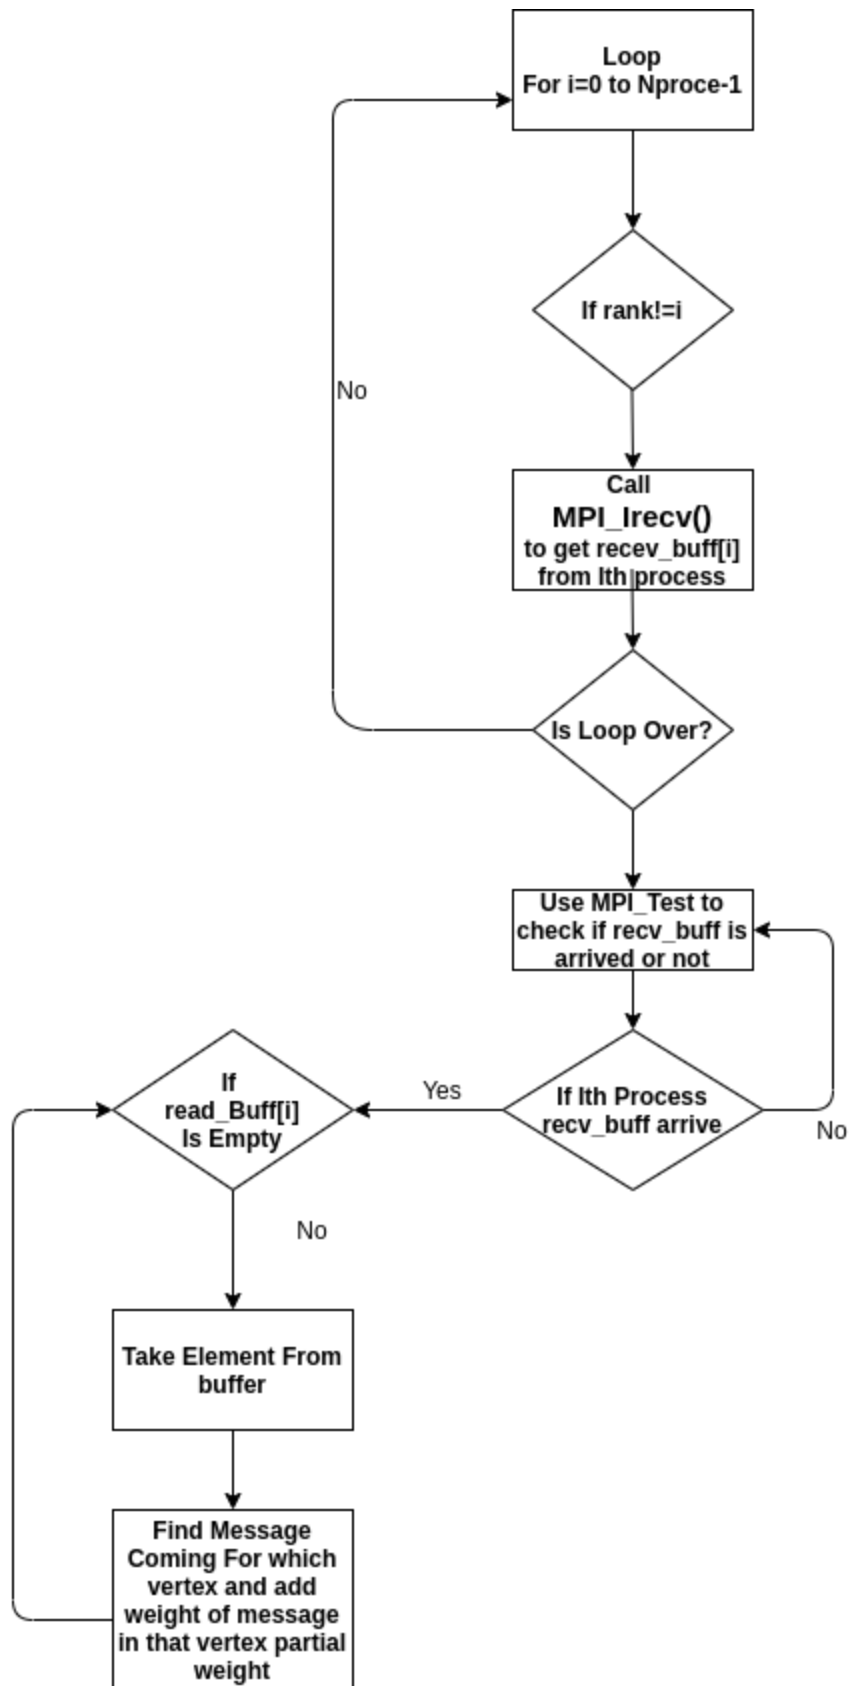
Using this approach I am getting a significant amount of performance improvement and speed up for the large graph.

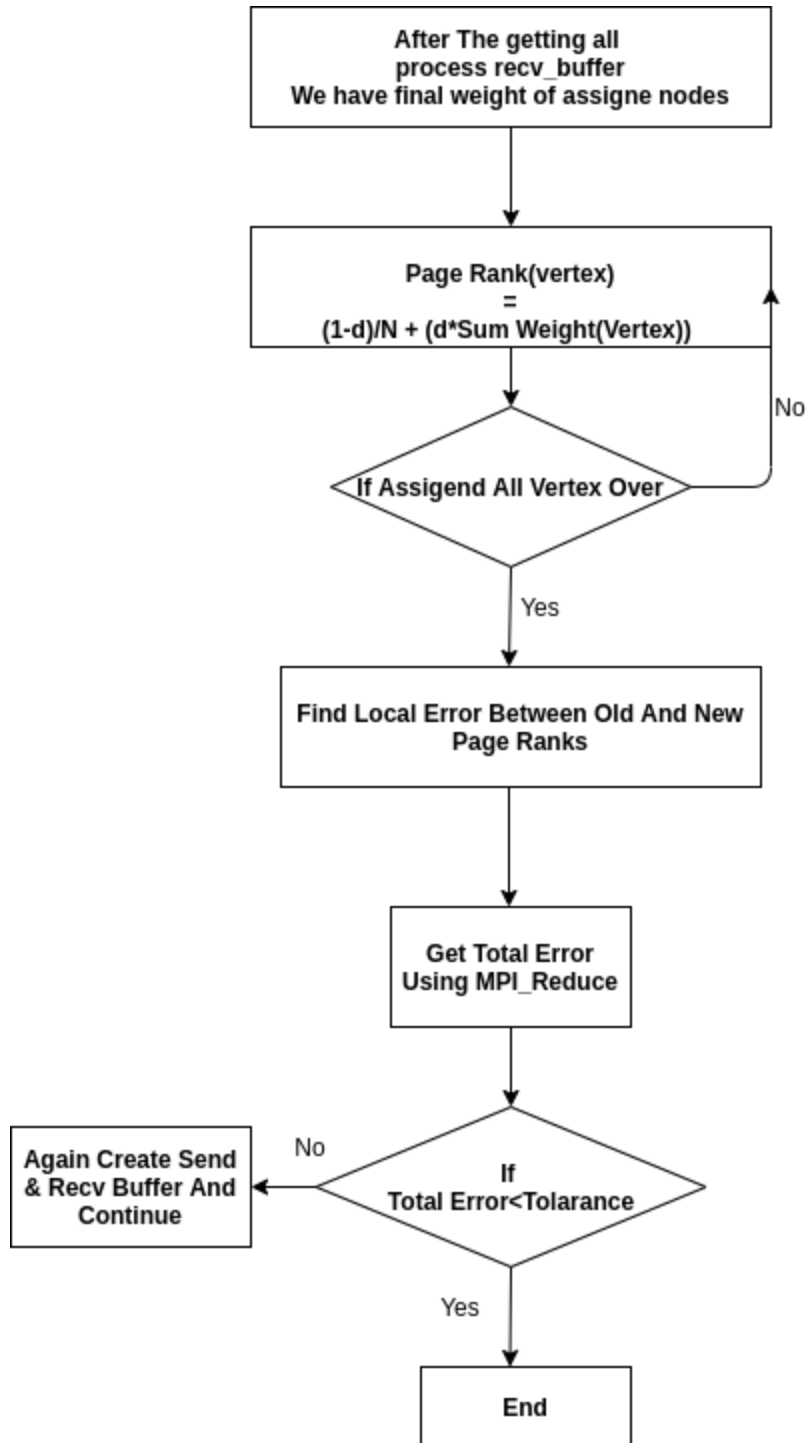
Process :

Procedure Followed By Each Process







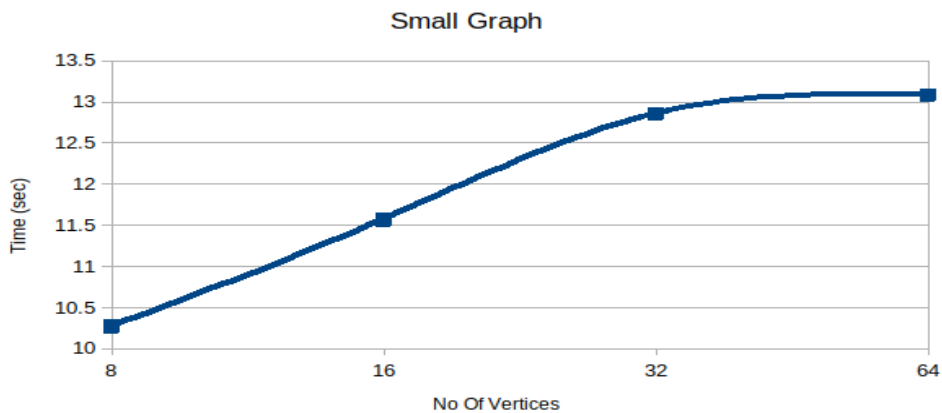


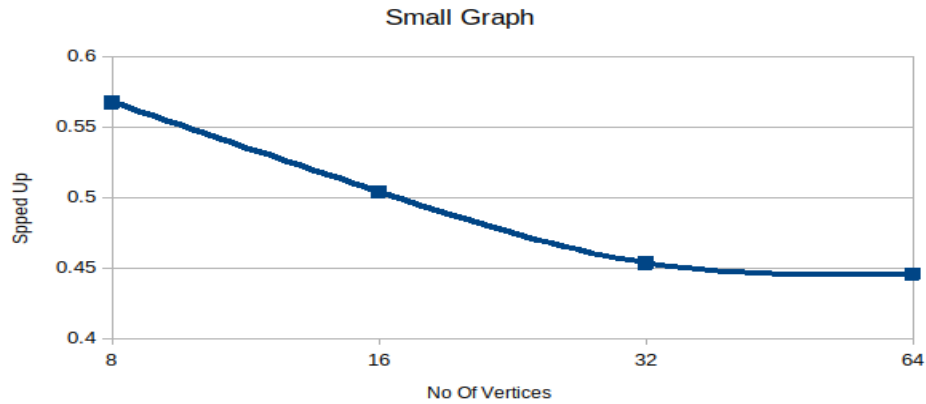
Performance Of The Last Approach :

Small Graph :

Sequential Execution Time	
Execution Number	Time (Sec)
1	5.86435
2	5.89143
3	5.87541
4	5.76951
5	5.73856
Average Execution Time	5.827852

No Of Processes	1	2	3	4	5	Average Time(sec)	Speed Up
8	10.2349	9.87571	10.382	10.3375	10.519	10.269822	0.567473516
16	10.9693	11.6272	11.9466	11.6945	11.6042	11.56836	0.50377512
32	12.2061	13.388	13.0413	13.451	12.1792	12.85312	0.45341924
64	12.8356	12.6631	13.3358	13.5801	12.9821	13.07934	0.44557691

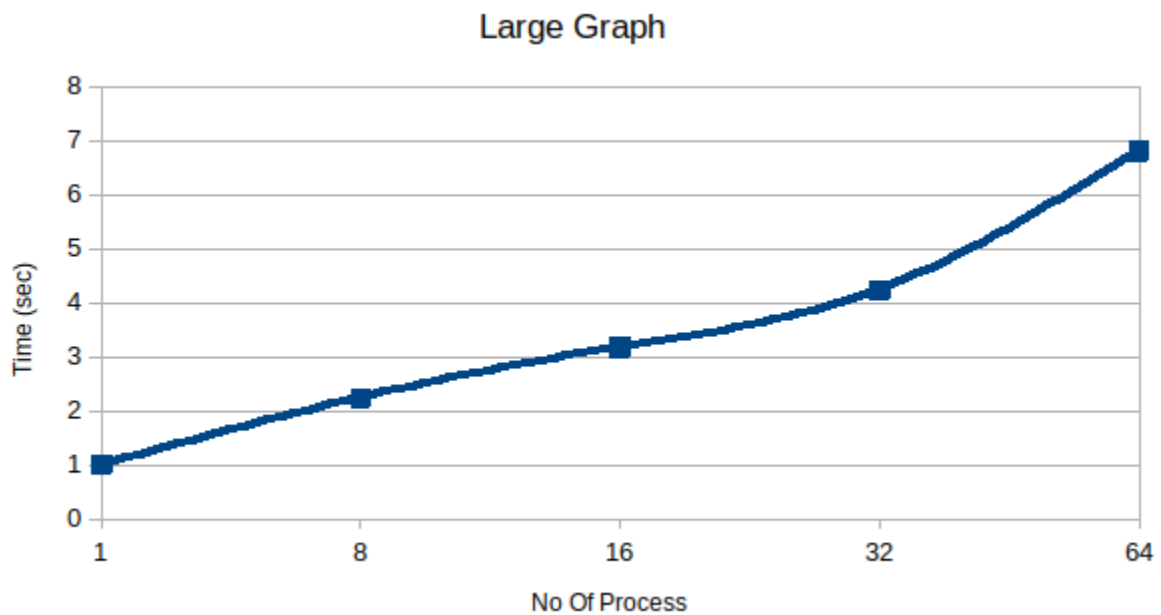
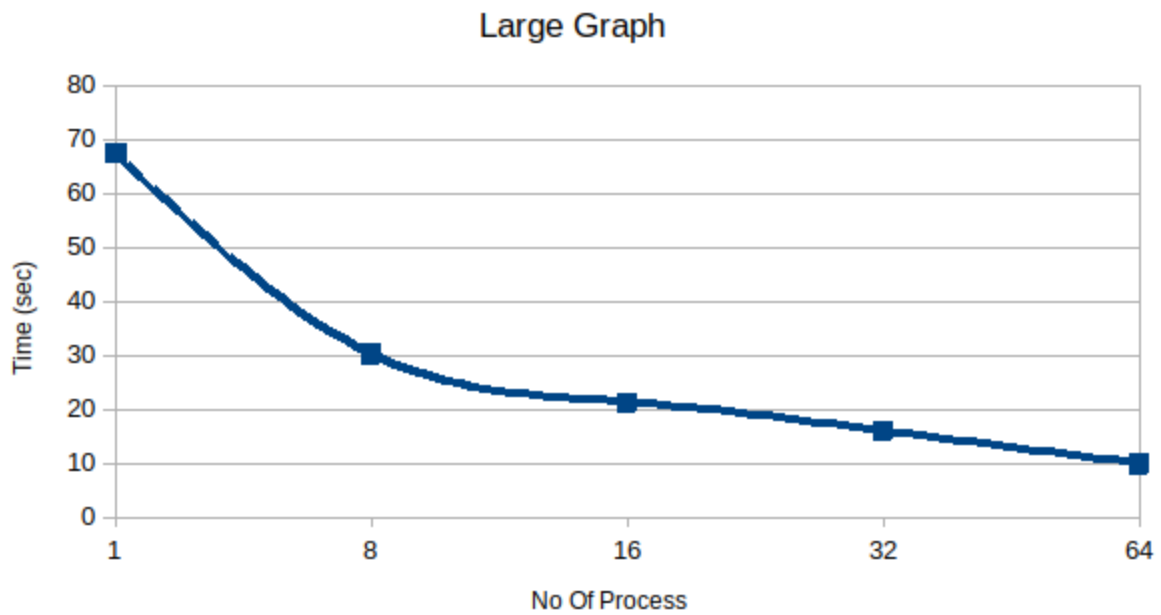




Large Graph

Sequential Execution Time	
Execution Number	Time (Sec)
1	64.8979
2	74.6496
3	60.5902
4	63.337
5	73.7916
Average Execution Time	67.45326

No Of Processes	1	2	3	4	5	Average Time(sec)	Speed Up
8	30.2417	29.9857	30.2145	30.3854	30.1478	30.19502	2.2339200
16	20.4125	21.5214	20.9821	21.6752	21.6542	21.249086	3.1744075
32	16.239	16.0176	15.9898	15.2555	16.1792	15.93622	4.2327013
64	9.59988	9.55722	10.1815	10.288	9.8965	9.90462	6.8102824



Remarks :

We can still get better performance if we will send only weight which has changed in the previous iteration.